

Notebook

March 15, 2025

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
```

```
[4]: # Read data
data = pd.read_csv('Learning_Intern_Raw_Data.csv', sep=',')
data
```

```
[4]:
```

	user_id	gender	age_range	phone \
0	6531ed09-304e-9573-a9c2-54e917	FEMALE	25-29	["23409089964086"]
1	6531ed09-304e-9573-a9c2-54e917	FEMALE	25-29	["23409089964086"]
2	6531ed09-304e-9573-a9c2-54e917	FEMALE	25-29	["23409089964086"]
3	6531ed09-304e-9573-a9c2-54e917	FEMALE	25-29	["23409089964086"]
4	6531ed09-304e-9573-a9c2-54e917	FEMALE	25-29	["23409089964086"]
...
65528	95539e11-a210-2b56-d5a4-a1b85f	FEMALE	30-34	["23408067849026"]
65529	95539e11-a210-2b56-d5a4-a1b85f	FEMALE	30-34	["23408067849026"]
65530	95539e11-a210-2b56-d5a4-a1b85f	FEMALE	30-34	["23408067849026"]
65531	95539e11-a210-2b56-d5a4-a1b85f	FEMALE	30-34	["23408067849026"]
65532	95539e11-a210-2b56-d5a4-a1b85f	FEMALE	30-34	["23408067849026"]

	state	program	assignment_name \
0	Lagos	Cloud Computing	Milestone #1
1	Lagos	Cloud Computing	Weekly Test #1
2	Lagos	Cloud Computing	Milestone #2
3	Lagos	Cloud Computing	Weekly Test #2
4	Lagos	Cloud Computing	Milestone #3
...
65528	Lagos	AI Career Essentials	Milestone #5
65529	Lagos	AI Career Essentials	Weekly Test #5
65530	Lagos	AI Career Essentials	Milestone #6
65531	Lagos	AI Career Essentials	Weekly Test #6
65532	Lagos	AI Career Essentials	AiCE Final Grace Period Test

	assignment_type	assignment_score	has_logged_into_lms \
0	milestone	100.00	Yes

1	test	95.23	Yes
2	milestone	94.55	Yes
3	test	90.90	Yes
4	milestone	100.00	Yes
...
65528	milestone	100.00	Yes
65529	test	100.00	Yes
65530	milestone	100.00	Yes
65531	test	100.00	Yes
65532	test	0.00	Yes

	Is assignment resubmitted	learner_deferred	learner_dropped_off	\
0	No	No	NaN	
1	No	No	NaN	
2	No	No	NaN	
3	No	No	NaN	
4	No	No	NaN	
...	
65528	No	No	NaN	
65529	Yes	No	NaN	
65530	No	No	NaN	
65531	Yes	No	NaN	
65532	No	No	NaN	

	overall_score	graduated
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN
...
65528	NaN	NaN
65529	NaN	NaN
65530	NaN	NaN
65531	NaN	NaN
65532	NaN	NaN

[65533 rows x 15 columns]

```
[5]: # Drop unnecessary columns
df = data[['user_id', 'program', 'assignment_name', 'assignment_type',
↪ 'assignment_score', 'has_logged_into_lms', 'Is assignment resubmitted',
↪ 'learner_deferred', 'learner_dropped_off', 'overall_score', 'graduated']]
```

```
[13]: # Remove deferred users
df = df[df["learner_deferred"] == "No"]
```

```
print("Deferred users removed. Updated dataset saved.")
```

Deferred users removed. Updated dataset saved.

```
[6]: # Filter only 'test' assignment types for overall_score
test_scores = df[df['assignment_type'] == 'test']
test_scores
```

```
[6]:
```

	user_id	program \
1	6531ed09-304e-9573-a9c2-54e917	Cloud Computing
3	6531ed09-304e-9573-a9c2-54e917	Cloud Computing
5	6531ed09-304e-9573-a9c2-54e917	Cloud Computing
7	6531ed09-304e-9573-a9c2-54e917	Cloud Computing
9	6531ed09-304e-9573-a9c2-54e917	Cloud Computing
...
65525	95539e11-a210-2b56-d5a4-a1b85f	AI Career Essentials
65527	95539e11-a210-2b56-d5a4-a1b85f	AI Career Essentials
65529	95539e11-a210-2b56-d5a4-a1b85f	AI Career Essentials
65531	95539e11-a210-2b56-d5a4-a1b85f	AI Career Essentials
65532	95539e11-a210-2b56-d5a4-a1b85f	AI Career Essentials

	assignment_name	assignment_type	assignment_score \
1	Weekly Test #1	test	95.23
3	Weekly Test #2	test	90.90
5	Weekly Test #3	test	100.00
7	Weekly Test #4	test	0.00
9	Weekly Test #5	test	0.00
...
65525	Weekly Test #3	test	100.00
65527	Weekly Test #4	test	80.95
65529	Weekly Test #5	test	100.00
65531	Weekly Test #6	test	100.00
65532	AiCE Final Grace Period Test	test	0.00

	has_logged_into_lms	Is assignment resubmitted	learner_deferred \
1	Yes	No	No
3	Yes	No	No
5	Yes	No	No
7	Yes	No	No
9	Yes	No	No
...
65525	Yes	No	No
65527	Yes	Yes	No
65529	Yes	Yes	No
65531	Yes	Yes	No
65532	Yes	No	No

	learner_dropped_off	overall_score	graduated
1	NaN	NaN	NaN
3	NaN	NaN	NaN
5	NaN	NaN	NaN
7	NaN	NaN	NaN
9	NaN	NaN	NaN
...
65525	NaN	NaN	NaN
65527	NaN	NaN	NaN
65529	NaN	NaN	NaN
65531	NaN	NaN	NaN
65532	NaN	NaN	NaN

[35287 rows x 11 columns]

```
[7]: # Compute average test score per user
overall_scores = test_scores.groupby('user_id')['assignment_score'].mean()

# Merge computed scores back into the original DataFrame
df.loc[:, 'overall_score'] = df['user_id'].map(overall_scores)

# Save the updated file
df.to_excel("Updated_Learning_Intern.xlsx", index=False)

print("Overall scores updated successfully!")
```

Overall scores updated successfully!

```
[18]: # Compute average overall score per user
avg_scores = df.groupby('user_id')['overall_score'].mean().reset_index()

# Create an interactive histogram
fig = px.histogram(avg_scores, x="overall_score", nbins=20,
                  title="Interactive Histogram: Distribution of Overall_
↪ Scores",
                  labels={"overall_score": "Average Overall Score", "count":
↪ "Number of Learners"},
                  opacity=0.75, color_discrete_sequence=["blue"])

# Customize hover info
fig.update_traces(marker_line_width=1, marker_line_color="black",
↪ hoverinfo="x+y")

# Show interactive plot
fig.show();
```

Over 2,600 Learners (more than half of the total) had test scores below 25. 1,901 learners achieved

test scores between 75 - 100

```
[24]: # Group by two columns and compute the mean overall score
grouped_df = df.groupby('program')['overall_score'].mean().reset_index()

# Create a bar chart
fig = px.bar(grouped_df, x="program", y="overall_score", color="program",
             title="Overall Test Score by Program",
             labels={"overall_score": "Average Test Score", "program": "Program"},
             barmode="group",
             color_discrete_sequence=px.colors.qualitative.Set1)

fig.show();
```

All programs have relatively the same average test scores.

```
[8]: df['overall_score']
```

```
[8]: 0      40.875714
     1      40.875714
     2      40.875714
     3      40.875714
     4      40.875714
     ...
    65528    81.692857
    65529    81.692857
    65530    81.692857
    65531    81.692857
    65532    81.692857
     Name: overall_score, Length: 65533, dtype: float64
```

```
[25]: # Define the required milestones for Learner_drop_off
required_milestones = ["Milestone #1", "Milestone #2", "Milestone #3"]

# Filter only required milestones
milestone_scores = df[df['assignment_name'].
                      isin(required_milestones)][['user_id', 'assignment_name',
                      'assignment_score']]

# Identify users who have **all three** milestones with a score greater than 0
valid_users = (
    milestone_scores.groupby('user_id') # Exclude 'user_id' from groups
    .apply(lambda x: all(x.set_index('assignment_name').
                        loc[required_milestones, 'assignment_score'] > 0))
    .reset_index()
)
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_14008\2721981694.py:10:
DeprecationWarning:

DataFrameGroupBy.apply operated on the grouping columns. This behavior is deprecated, and in a future version of pandas the grouping columns will be excluded from the operation. Either pass `include_groups=False` to exclude the groupings or explicitly select the grouping columns after groupby to silence this warning.

```
[27]: # Get user IDs of those who submitted required milestones
submitted_all = valid_users[valid_users[0]]['user_id']

# Assign learner_dropped_off: 1 (Dropped Off), 0 (Stayed)
df.loc[:, 'learner_dropped_off'] = df['user_id'].apply(lambda x: 0 if x in submitted_all.values else 1)

# Validation Step: Count dropped users
dropped_count = int(df['learner_dropped_off'].sum()/13)
total_users = df['user_id'].nunique()
print(f"Total Learners: {total_users}")
print(f"Learners Dropped Off: {dropped_count}")
print(f"Learners Remaining: {total_users - dropped_count}")

# Save the updated DataFrame
df.to_csv("updated_data.csv", index=False)

print("Users who didn't submit valid milestone scores have been marked as dropped off successfully!")
```

Total Learners: 5002

Learners Dropped Off: 3040

Learners Remaining: 1962


Users who didn't submit valid milestone scores have been marked as dropped off successfully!

```
[28]: # Calculate percentages
dropped_off_percentage = (dropped_count / total_users) * 100
remaining_percentage = ((total_users - dropped_count) / total_users) * 100

# Data for plotting
labels = ['Dropped Off', 'Remaining']
percentages = [dropped_off_percentage, remaining_percentage]

# Create a bar chart
plt.figure(figsize=(8, 6))
plt.bar(labels, percentages, color=['red', 'green'])
```

```
plt.xlabel('Learner Status')
plt.ylabel('Percentage')
plt.title('Percentage of Learners Dropped Off vs Remaining')
plt.ylim(0, 100) # Set y-axis limit to 100%
plt.show();
```



```
[29]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 65026 entries, 0 to 65532
Data columns (total 11 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   user_id                               65026 non-null  object
1   program                               65026 non-null  object
2   assignment_name                       65026 non-null  object
3   assignment_type                       65026 non-null  object
4   assignment_score                      65026 non-null  float64
5   has_logged_into_lms                  65026 non-null  object
6   Is assignment resubmitted             65026 non-null  object
```

```

7 learner_deferred          65026 non-null object
8 learner_dropped_off      65026 non-null float64
9 overall_score            65026 non-null float64
10 graduated                0 non-null float64
dtypes: float64(4), object(7)
memory usage: 6.0+ MB

```

```

[39]: # Define the required milestones to graduate
reqd_milestones = ["Milestone #1", "Milestone #2", "Milestone #3", "Milestone #4", "Milestone #5"]

# Filter only required milestones
mile_scores = df[df['assignment_name'].isin(reqd_milestones)][['user_id', 'assignment_name', 'assignment_score']]

# Drop rows where assignment_score is 0.00
mile_scores = mile_scores[mile_scores["assignment_score"] > 0]

```

```

[40]: # Users who submitted all 5 milestones
f_mile_scores = mile_scores.groupby('user_id')['assignment_name'].count() > 4

# Filter the dataset to retain only 'f_mile_scores'
filtered_mile_scores = mile_scores[mile_scores['user_id'].isin(f_mile_scores[f_mile_scores].index)]

# Display result
filtered_mile_scores

```

```

[40]:
      user_id assignment_name assignment_score
39  02f0e023-398a-93e5-217e-5448f0  Milestone #1          96.55
41  02f0e023-398a-93e5-217e-5448f0  Milestone #2         100.00
43  02f0e023-398a-93e5-217e-5448f0  Milestone #3         100.00
45  02f0e023-398a-93e5-217e-5448f0  Milestone #4          99.68
47  02f0e023-398a-93e5-217e-5448f0  Milestone #5          96.67
...
65520  95539e11-a210-2b56-d5a4-a1b85f  Milestone #1         100.00
65522  95539e11-a210-2b56-d5a4-a1b85f  Milestone #2          99.32
65524  95539e11-a210-2b56-d5a4-a1b85f  Milestone #3          98.53
65526  95539e11-a210-2b56-d5a4-a1b85f  Milestone #4         100.00
65528  95539e11-a210-2b56-d5a4-a1b85f  Milestone #5         100.00

```

[8670 rows x 3 columns]

```
[ ]: f_mile_scores.info()
```

```

<class 'pandas.core.series.Series'>
Index: 2540 entries, 0034a903-0652-77b3-9c70-82b91a to

```



```

fff94200-23ac-3b93-fca0-178543
Series name: assignment_name
Non-Null Count  Dtype
-----
2540 non-null   bool
dtypes: bool(1)
memory usage: 86.9+ KB

```

```

[41]: # Method 2 to obtain graduation milestone requirement

# Pivot table to ensure each user has all 5 milestones
milestone_pivot = mile_scores.pivot(index="user_id", columns="assignment_name",
    ↪values="assignment_score")

# Check if each user has ALL Milestones (#1 to #5)
users_with_5_milestones = milestone_pivot[reqd_milestones].gt(0).all(axis=1)

# Output Results
users_with_5_milestones

```

```

[41]: user_id
0034a903-0652-77b3-9c70-82b91a    False
005b9685-1114-e183-1b69-45964c    False
005d3e80-f153-c133-8dea-4e5b66     True
0065897d-5ddf-d164-00c2-550e2e    False
00a5a86d-343d-3a27-e9ab-be6f4f    False
...
ffbaaeb2-ab3c-03d3-a765-b5d611    True
ffd40f70-6761-dabd-b0b7-2c486c    False
ffd68c7b-8edc-9e0b-45f6-168767    True
fff8b323-3e76-e67d-cf1f-8390fb    False
fff94200-23ac-3b93-fca0-178543    True
Length: 2540, dtype: bool

```

```

[ ]: milestone_pivot.head()

```

```

[ ]: assignment_name      Milestone #1  Milestone #2  Milestone #3  \
user_id
0034a903-0652-77b3-9c70-82b91a      100.00         NaN         NaN
005b9685-1114-e183-1b69-45964c         NaN        87.11        27.27
005d3e80-f153-c133-8dea-4e5b66        88.97        98.64        91.18
0065897d-5ddf-d164-00c2-550e2e      100.00         NaN         NaN
00a5a86d-343d-3a27-e9ab-be6f4f         NaN         NaN        83.53
00b795d6-f7b8-7904-f685-bef7d7        89.66        91.59       100.00
00d56c28-1c61-e176-fbd4-6acb6c        89.31        81.82        66.67
00f46499-ad02-18ae-e634-9b5613        87.93        94.77        87.94
00fd6e8f-8ace-3e32-1b47-d9d0ed      100.00        98.64       100.00

```

012a5010-4041-71ac-c5f3-ec0ced	100.00	53.64	100.00
--------------------------------	--------	-------	--------

assignment_name	Milestone #4	Milestone #5
user_id		
0034a903-0652-77b3-9c70-82b91a	NaN	NaN
005b9685-1114-e183-1b69-45964c	NaN	NaN
005d3e80-f153-c133-8dea-4e5b66	100.00	93.70
0065897d-5ddf-d164-00c2-550e2e	NaN	NaN
00a5a86d-343d-3a27-e9ab-be6f4f	70.97	61.85
00b795d6-f7b8-7904-f685-bef7d7	94.19	100.00
00d56c28-1c61-e176-fbd4-6acb6c	77.42	98.89
00f46499-ad02-18ae-e634-9b5613	95.16	90.37
00fd6e8f-8ace-3e32-1b47-d9d0ed	96.77	96.30
012a5010-4041-71ac-c5f3-ec0ced	90.00	97.04

```
[42]: # Extract test data required for graduation
tests = df[df['assignment_name'].str.contains("Weekly Test #", na=False)]

# Ensure users have an average test score 75%
test_scores = tests.pivot(index='user_id', columns='assignment_name',
    ↪values='assignment_score').fillna(0)
test_avg = test_scores.mean(axis=1) # Compute average score
test_criteria = test_avg >= 75 # True if average test score is at least 75%
```

```
[ ]: test_criteria
```

```
[ ]: user_id
0034a903-0652-77b3-9c70-82b91a    False
005a2866-1063-8f09-bf83-1cb6f4    False
005b9685-1114-e183-1b69-45964c    False
005d3e80-f153-c133-8dea-4e5b66     True
005f8882-251f-a015-ff42-c8e80c    False
...
ffc80062-9d1b-3d92-6647-d46d08    False
ffd40f70-6761-dabd-b0b7-2c486c    False
ffd68c7b-8edc-9e0b-45f6-168767     True
fff8b323-3e76-e67d-cf1f-8390fb    False
fff94200-23ac-3b93-fca0-178543     True
Length: 5041, dtype: bool
```

```
[43]: # Apply both conditions to determine valid graduates
graduation_status = (users_with_5_milestones & test_criteria).astype(int)

# Assign graduation status to the original DataFrame
df['graduated'] = df['user_id'].map(graduation_status).fillna(0).astype(int)

print("Graduation status updated.")
```

Graduation status updated.

```
[44]: # Count unique graduates
num_graduated = df["user_id"][df["graduated"] == 1].nunique()
print(f"Total number of users who graduated: {num_graduated}")
```


Total number of users who graduated: 1712

```
[30]: # Calculate percentages
total_learners = 5002
dropped_off = 3290
graduated = 1712

dropped_off_percentage = (dropped_off / total_learners) * 100
graduated_percentage = (graduated / total_learners) * 100

# Data for plotting
labels = ['Graduated', 'Dropped Off']
percentages = [graduated_percentage, dropped_off_percentage]

# Create a bar chart
plt.figure(figsize=(8, 6))
plt.bar(labels, percentages, color=['red', 'green'])
plt.xlabel('Learner Status')
plt.ylabel('Percentage')
plt.title('Percentage of Learners Dropped Off vs Graduated')
plt.ylim(0, 100) # Set y-axis limit to 100%
plt.show();
```



Less than 40% of the total learners graduated

```
[ ]: # Method 2 to identify valid graduates
test_criteria_df = test_criteria.to_frame(name='test_passed') # Test criteria
f_mile_scores_df = f_mile_scores.to_frame(name='milestone_completed') #
↳ Milestone criteria

# Merge both graduation criteria
final_df = pd.merge(test_criteria_df, f_mile_scores_df, left_index=True,
↳ right_index=True, how='inner')
```

```
[ ]: final_df
```

```
[ ]:
          test_passed  milestone_completed
user_id
0034a903-0652-77b3-9c70-82b91a      False      False
005b9685-1114-e183-1b69-45964c      False      False
005d3e80-f153-c133-8dea-4e5b66         True         True
0065897d-5ddf-d164-00c2-550e2e      False      False
00a5a86d-343d-3a27-e9ab-be6f4f         True      False
```

...
ffbaaeb2-ab3c-03d3-a765-b5d611	True	True
ffd40f70-6761-dabd-b0b7-2c486c	False	False
ffd68c7b-8edc-9e0b-45f6-168767	True	True
fff8b323-3e76-e67d-cf1f-8390fb	False	False
fff94200-23ac-3b93-fca0-178543	True	True

[2540 rows x 2 columns]

```
[ ]: # Check for learners who meet both conditions
final_df = final_df[(final_df['test_passed'] == True) &
    ↪(final_df['milestone_completed'] == True) ]
```

```
[ ]: # Obtain user_id of graduating learners
f_user_id = final_df.index
f_user_id
```

```
[ ]: Index(['005d3e80-f153-c133-8dea-4e5b66', '00b795d6-f7b8-7904-f685-bef7d7',
          '00d56c28-1c61-e176-fbd4-6acbac', '00f46499-ad02-18ae-e634-9b5613',
          '00fd6e8f-8ace-3e32-1b47-d9d0ed', '012a5010-4041-71ac-c5f3-ec0ced',
          '014c1879-2025-1109-297f-57e7d8', '01695c10-c70c-9ad3-c211-573a88',
          '019fe0fc-c55c-1481-1170-ea7bb2', '01c6cef4-187b-9d11-efa8-d6fdbb',
          ...
          'fed0ceb6-f5f7-273a-ba35-7fc573', 'fefa6261-1ce8-f1bb-71ef-ae8ca0',
          'ff2d5b94-dab8-0977-a4dc-c92ecd', 'ff470390-9fad-adbd-40b3-391bb3',
          'ff4e42e6-d026-25b6-cecf-5471b9', 'ff88bf17-7269-e3ad-cfd5-a4a622',
          'ffb0ecf0-eb0d-85c0-0d8b-aac797', 'ffbaaeb2-ab3c-03d3-a765-b5d611',
          'ffd68c7b-8edc-9e0b-45f6-168767', 'fff94200-23ac-3b93-fca0-178543'],
          dtype='object', name='user_id', length=1712)
```

```
[ ]: # Populate 'graduate' column based on user_id present in f_user_id
df['graduate'] = df['user_id'].isin(f_user_id).astype(int)

# Count unique graduates
unique_graduates = df["user_id"][df["graduate"] == 1].nunique()
print(f"Total number of unique graduates: {unique_graduates}")
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_11800\2048875875.py:6:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df['graduate'] = df['user_id'].isin(f_user_id).astype(int)
```

```
[ ]: # Validation output for both methods used
mismatch_count = (df['graduated'] != df['graduate']).sum()
print(mismatch_count)
```

0

Both methods yielded the exact result for graduating learners.

```
[50]: # Group by program and compute graduation rate
graduation_rates = df.groupby('program')['graduated'].mean() * 100 # Convert_
    ↳ to percentage

# Sort values for better visualization
graduation_rates = graduation_rates.sort_values(ascending=False)

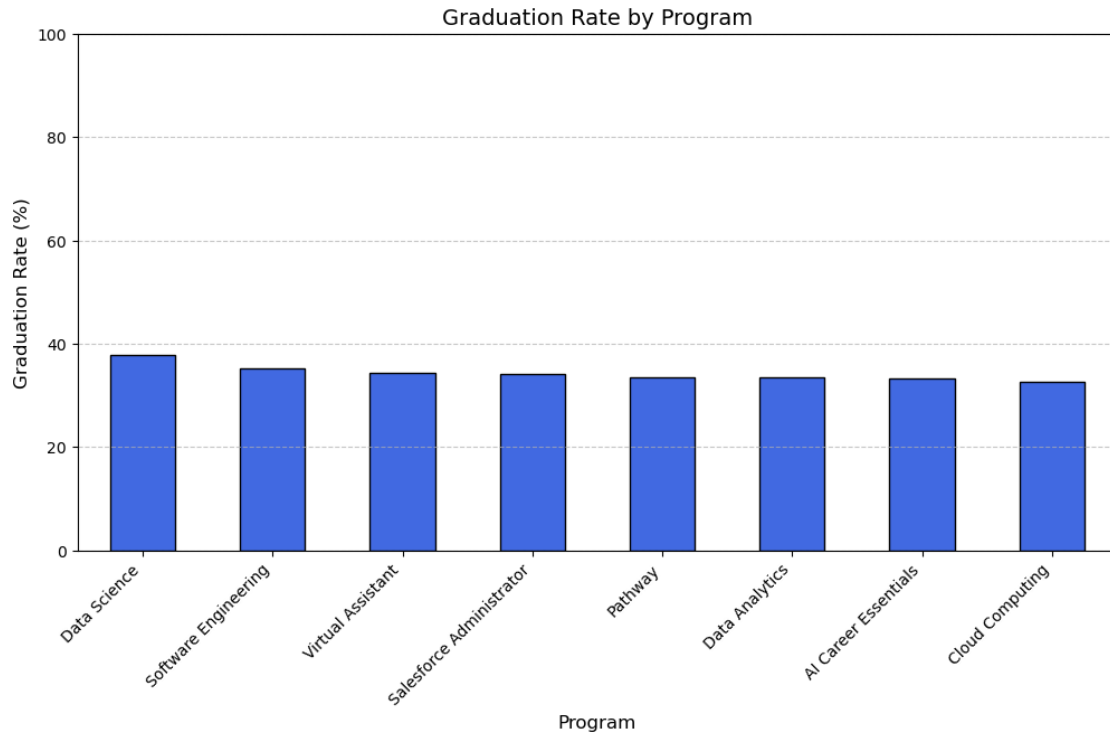
print(graduation_rates) # Check the computed graduation rates
```

```
program
Data Science          37.730871
Software Engineering  35.270806
Virtual Assistant     34.274953
Salesforce Administrator 34.239130
Pathway               33.522727
Data Analytics        33.429672
AI Career Essentials  33.333333
Cloud Computing       32.647059
Name: graduated, dtype: float64
```

```
[51]: plt.figure(figsize=(12, 6)) # Set figure size
graduation_rates.plot(kind='bar', color='royalblue', edgecolor='black')

plt.title("Graduation Rate by Program", fontsize=14)
plt.xlabel("Program", fontsize=12)
plt.ylabel("Graduation Rate (%)", fontsize=12)
plt.xticks(rotation=45, ha='right') # Rotate x-axis labels for readability
plt.ylim(0, 100) # Ensure the y-axis shows percentages properly

plt.grid(axis='y', linestyle='--', alpha=0.7) # Add gridlines for better_
    ↳ readability
plt.show();
```



The graduation rate by program chart reveals that the program is not a